

DATA-236 Sec 12 - Distributed Systems for Data Engineering
HOMEWORK 9
Nandhakumar Apparsamy
018190003

GitHub - <https://github.com/Nandha951/DATA-236-HW9-Redis-FastAPI-MongoDB>

Make a student schema consisting of the following fields

- Student Name (String)
- Study_hours (Number)
- Attendance (Number)
- Score (Number)

```
const studentSchema = new mongoose.Schema({
  Student_Name: String,
  Study_hours: Number,
  Attendance: Number,
  Score: Number
});
```

Redis Setup Instructions

1. Pull the official redis image from Docker Hub

docker pull redis

```
● (base) spartan@MLK-SCS-P0WGL9N2QF HW9 % docker pull redis
Using default tag: latest
latest: Pulling from library/redis
Digest: sha256:fbdbaea47b9ae4ecc2082ecdb4e1cea81e32176ffb1dcf643d422ad07427e5d9
Status: Image is up to date for redis:latest
docker.io/library/redis:latest
○ (base) spartan@MLK-SCS-P0WGL9N2QF HW9 %
```

2.Run the redis container

`docker run -d --name redis-server -p 6379:6379 redis`

```
Run 'docker run --help' for more information
● (base) spartan@MLK-SCS-P0WGL9N2QF HW9 % docker run -d --name redis-server -p 6379
:6379 redis
3a8d87af7103bc7c40b8ede01ad3f1d723e9316e2a2d53184999c206bd04dcaf
○ (base) spartan@MLK-SCS-P0WGL9N2QF HW9 %
```

Command to Start FastAPI

`uvicorn app:app --host 0.0.0.0 --port 8000 --reload`

```
○ (base) spartan@MLK-SCS-P0WGL9N2QF fastapi_app % uvicorn main:app --host 0.0.0.0 --
-port 8000 --reload
INFO:      Will watch for changes in these directories: ['/Users/spartan/SJSU/DATA
236/HW Assignment/HW9/fastapi_app']
INFO:      Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
INFO:      Started reloader process [24302] using StatReload
INFO:      Started server process [24304]
INFO:      Waiting for application startup.
INFO:      Application startup complete.
```

Q1. Build a GET API using Node.js to fetch all student records from MongoDB(4 Points)

Container CPU usage ⓘ

0.24% / 800% (8 CPUs available)

Container memory usage ⓘ

9.49MB / 7.47GB

Show charts

🔍 Search

☰

Only show running containers

<input type="checkbox"/>	Name	Container ID	Image	Port(s)	CPU (%)	Las	Actions
<input type="checkbox"/>	<div><div></div>redis-server</div>	3a8d87af7103	redis	6379:6379 ↗	0.48%	4 m	<div><div><div></div></div><div><div></div></div><div><div></div></div></div>

- Store the fetched result in Redis.
- (Refer: <https://redis.io/docs/latest/develop/clients/nodejs/>)
- If the data exists in Redis, return it from the cache.
- Log where the data is being served from (Redis or MongoDB).

```
// Redis Connection
const redisClient = redis.createClient({
  host: 'redis-server',
  port: 6379
});

redisClient.on('connect', () => console.log('Connected to Redis'));
redisClient.on('error', err => console.error('Redis connection error:', err))

redisClient.connect().catch(console.error)

// Express middleware to parse JSON
app.use(express.json());

// Function to fetch student records from MongoDB
async function fetchStudentRecords() {
  try {
    const students = await Student.find();
    console.log('Data served from MongoDB:', students);
    return students;
  } catch (error) {
    console.error('Error fetching student records from MongoDB:', error);
    throw error;
  }
}
```

```
// Function to store student records in Redis
async function storeStudentRecordsInRedis(students) {
  try {
    await redisClient.set('students', JSON.stringify(students));
    console.log('Student records stored in Redis');
  } catch (error) {
    console.error('Error storing student records in Redis:', error);
    throw error;
  }
}

// Function to get student records from Redis
async function getStudentRecordsFromRedis() {
  try {
    const students = await redisClient.get('students');
    if (students) {
      console.log('Data served from Redis:', students);
      return JSON.parse(students);
    }
    return null;
  } catch (error) {
    console.error('Error getting student records from Redis:', error);
    return null;
  }
}
```

```
// GET API to fetch all student records
app.get('/students', async (req, res) => {
  try {
    let students = await getStudentRecordsFromRedis();
    if (!students) {
      students = await fetchStudentRecords();
      await storeStudentRecordsInRedis(students);
    }
    res.json(students);
  } catch (error) {
    console.error('Error fetching student records:', error);
    res.status(500).json({ error: 'Failed to fetch student records' });
  }
});
```

```
Node.js app listening at http://localhost:3000
Connected to Redis
Connected to MongoDB
Data already exists in MongoDB, skipping CSV load.
Data served from Redis: [{"_id":"67f81d4906ec95eaa8a09433","Student_Name":"Updated Name","Study_hours":4.4,"Attendance":84.4,"Score":70,"__v":0}, {"_id":"67f81d4906ec95eaa8a09434","Student_Name":"Unknown","Study_hours":9.6,"Attendance":66.8,"Score":122,"__v":0}, {"_id":"67f81d4906ec95eaa8a09435","Student_Name":"Unknown","Study_hours":7.6,"Attendance":62.6,"Score":103,"__v":0}, {"_id":"67f81d4906ec95eaa8a09436","Student_Name":"Unknown","Study_hours":6.4,"Attendance":98,"Score":124,"__v":0}, {"_id":"67f81d4906ec95eaa8a09437","Student_Name":"Unknown","Study_hours":2.4
```



Student API / **Get All Students**

GET



http://localhost:3000/students

Params

Authorization

Headers (6)

Body

Scripts

Settings

Query Params

	Key	Value
	Key	Value

Body

Cookies

Headers (7)

Test Results



200 OK

{} JSON



▶ Preview

🔄 Visualize




```
1  [  
2    {  
3      "_id": "67f81d4906ec95eaa8a09433",  
4      "Student_Name": "Updated Name",  
5      "Study_hours": 4.4,  
6      "Attendance": 84.4,  
7      "Score": 70,  
8      "__v": 0  
9    },  
10   {  
11     "_id": "67f81d4906ec95eaa8a09434",  
12     "Student_Name": "Unknown",  
13     "Study_hours": 9.6,  
14     "Attendance": 66.8,  
15     "Score": 122,  
16     "__v": 0  
17   },  
18 ]
```

Q2. Use GitHub Copilot to assist in writing the PUT API to update a student's record.(2 Points)



Nandha951

Write code to develop PUT API to update a student' record

 Codebase



GitHub Copilot

✓ Searching codebase for "PUT API to update a student's record"...

```
// PUT API to update a student's record
app.put('/students/:id', async (req, res) => {
  try {
    const { id } = req.params;
    const { Student_Name, Study_hours, Attendance, Score } = req.body;

    // Find the student by ID and update their record
    const updatedStudent = await Student.findByIdAndUpdate(
      id,
      { Student_Name, Study_hours, Attendance, Score },
      { new: true } // Return the updated document
    );

    if (!updatedStudent) {
      return res.status(404).json({ error: 'Student not found' });
    }

    // Clear the Redis cache to ensure the next GET request fetches updated data
    await redisClient.del('students');

    res.json(updatedStudent);
  } catch (error) {
    console.error('Error updating student record:', error);
    res.status(500).json({ error: 'Failed to update student record' });
  }
});
```

HTTP Student API / Update Student

PUT http://localhost:3000/students/:id

Params Authorization Headers (9) Body Scripts Settings

☐ none ☐ form-data ☐ x-www-form-urlencoded ☒ raw ☐ binary ☐ G

```
1 {
2   "Student_Name": "Updated Name",
3   "Study_hours": 4.4,
4   "Attendance": 84.4,
5   "Score": 70
6 }
```

Body Cookies Headers (7) Test Results 200 OK

{ } JSON Preview Visualize

```
1 {
2   "_id": "67f81d4906ec95eaa8a09433",
3   "Student_Name": "Updated Name",
4   "Study_hours": 4.4,
5   "Attendance": 84.4,
6   "Score": 70,
7   "__v": 0
8 }
```

Q3. Build a score prediction system using FastAPI and a machine learning model.(4 Points)

- Use the provided dataset: Student_performance_dataset.csv
- Train a simple regression model to predict score based on study_hours and attendance.
- Save the model as a .pkl file.


```

import pandas as pd
from sklearn.linear_model import LinearRegression
import pickle
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel

# Load the dataset
data = None
try:
    data = pd.read_csv("Student_performance_dataset.csv")
except FileNotFoundError:
    raise FileNotFoundError("Student_performance_dataset.csv not found. Pleas

# Train the model
model = LinearRegression()
model.fit(data[['StudyHours', 'Attendance']], data['Score'])

# Save the model
filename = "score_prediction_model.pkl"
pickle.dump(model, open(filename, 'wb'))

# Load the model
loaded_model = pickle.load(open(filename, 'rb'))


```

- Create a FastAPI endpoint (/predict) that:
- Accepts study hours and attendance
- Returns the predicted score

```
# Create the FastAPI app
app = FastAPI()

# Define the input data model
class PredictionInput(BaseModel):
    StudyHours: float
    Attendance: float

# Create the prediction endpoint
@app.post("/predict")
async def predict_score(input_data: PredictionInput):
    try:
        study_hours = input_data.StudyHours
        attendance = input_data.Attendance
        print(f"Study hours: {study_hours}, Attendance: {attendance}")
        prediction = loaded_model.predict([[study_hours, attendance]])[0]
        return {"predicted_score": round(prediction, 2)}
    except Exception as e:
        print(f"Error: {e}")
        raise HTTPException(status_code=500, detail=str(e))

if __name__ == "__main__":
    # You can use uvicorn to run the app locally
    # uvicorn main:app --reload
     pass
```

- From your Node.js app, make a POST request to this FastAPI endpoint and return the prediction.

```
const fastAPIEndpoint = 'http://localhost:8000/predict';

// POST API to predict score
app.post('/predict', async (req, res) => {
  try {
    let StudyHours = req.body.StudyHours;
    let Attendance = req.body.Attendance;
    if (StudyHours === undefined || Attendance === undefined) {
      return res.status(400).json({ error: 'Missing StudyHours or Attendance' });
    }
    StudyHours = parseFloat(StudyHours);
    Attendance = parseFloat(Attendance);
    console.log(`Sending to FastAPI: StudyHours=${StudyHours}, Attendance=${Attendance}`);
    const response = await axios.post(fastAPIEndpoint, { StudyHours, Attendance });
    res.json(response.data);
  } catch (error) {
    console.error('Error predicting score:', error);
    res.status(500).json({ error: 'Failed to predict score' });
  }
});
```